

IMPLEMENTATION OF SIGNAL HANDLING

CS124 – Operating Systems
Fall 2018-2019, Lecture 15

Signal Handling

- UNIX operating systems allow user processes to register for and handle **signals**
 - Provides exceptional control flow mechanism for user processes
- User program registers a **signal handler** via a system call
- Example:

```
typedef void (*sig_t) (int);
```

 - `sig_t` is a function-pointer to a function that takes an `int` argument and returns `void`

```
sig_t signal(int sig, sig_t func);
```

 - This system call sets the signal handler for the specified signal type, and returns the previous signal handler
- Declarations are in C standard header `signal.h`

Signal Handling: Example

```
/* Print a message, then request another SIGALRM. */
void handle_sigalrm(int sig) {
    printf("Hello!\n");
    alarm(1); /* Request another SIGALRM in 1 second. */
}

/* User typed Ctrl-C. Taunt them. */
void handle_sigint(int sig) {
    printf("Ha ha, can't kill me!\n");
}

int main() {
    signal(SIGINT, handle_sigint);
    signal(SIGALRM, handle_sigalrm);
    alarm(1); /* Request a SIGALRM in 1 second. */

    while (1) pause(); /* Wait for signals in a loop. */

    return 0;
}
```

Advanced Signal Handling Support

- UNIX also provides more advanced signal handling:

```
int sigaction(int sig, const struct sigaction *act,  
             struct sigaction *oact)
```

- **sigaction** struct specifies various details, including the kind of handler function:

- Either the simple **void handler(int sig)** as before...
 - Or, a more advanced handler function:

```
void sigact(int sig, siginfo *info, void *ctxt)
```

- The **siginfo** struct includes many details about signals

- e.g. sending process ID, memory address that caused fault, etc.

- **ctxt** points to a **ucontext_t** structure

- A platform/architecture-dependent machine context, containing the CPU state of the user process, when it was interrupted by signal
 - Facilitates e.g. user-space threading libraries

Pending and Blocked Signals

- The kernel maintains two bit-vectors for every process
 - Every type of signal has a specific bit in this bit-vector
- **pending** bit-vector records what signals have yet to be delivered to the process
 - Note: if multiple instances of a given signal occur before a process receives the signal, it will see only one instance of the signal
 - Signals indicate one or more events of given type have occurred
- **blocked** bit-vector records what signals are currently not allowed to be delivered to the process
 - Can have a signal that is both blocked and pending
 - When the signal is unblocked, it will be delivered to the process
- When a signal is delivered to a process, that type of signal is automatically blocked for the process
 - Prevents a given signal handler from interrupting itself
 - One kind of signal can interrupt another kind of signal

The Kernel and Signal Handling

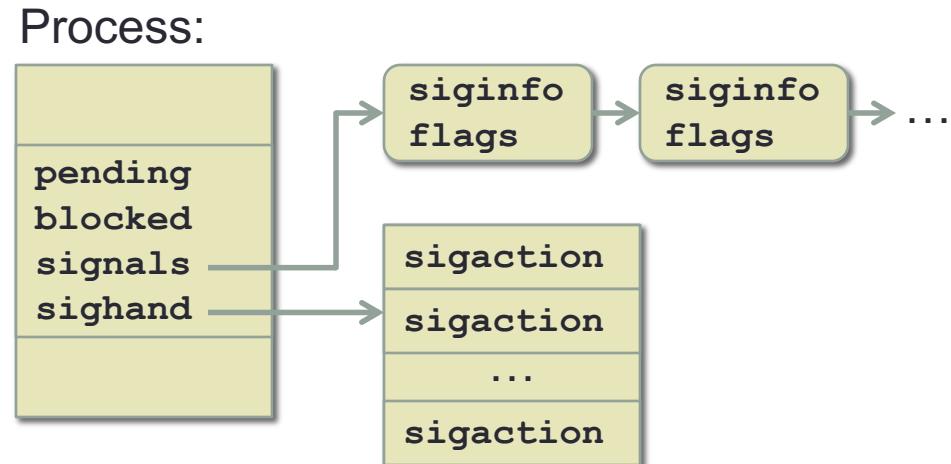
- How does the kernel provide signal handling to user processes?
- When is a signal delivered to the receiving process?
 - What if process is running on the CPU? What if ready or blocked?
 - How does signal handling affect process scheduling?
- If a process has multiple signals pending, how does the kernel dispatch all these signals to the process?
- Note: we are ignoring many issues caused by using signals in multithreaded programs
 - Individual threads can block signals so that only one thread handles signals, etc.

Generating and Delivering Signals

- A process isn't always running when a signal is sent to it
 - e.g. `kill()` syscall is invoked by another process
 - e.g. a child process dies, causing `SIGCHLD` to be sent to parent, but a higher priority process currently preempts the parent
- Kernels make a distinction between *generating* a signal and *delivering* the signal
- Signal generation: kernel updates the data structures of the receiving process to record that the signal was sent
- Signal delivery: kernel forces the receiving process to respond to the signal (e.g. by invoking a signal handler)
- Time may pass between generating and delivering signal

Process Signal Data Structures

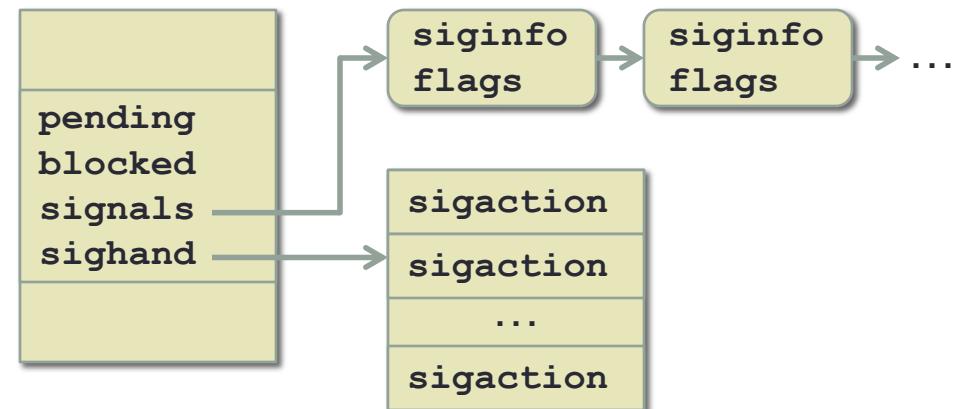
- Already mentioned **pending** / **blocked** signal bit-vectors
 - At a coarse-grain level of detail, records which signals need to be delivered, and which signals are currently blocked from delivery
- Each process also has a linked list of pending signals
 - `siginfo_t` struct records relevant details of the pending signal
- Each process also has an array of “signal action” structs
 - Specifies how to handle each kind of signal
 - e.g. “default action,” “ignore,” or a user-space handler
 - (Flags also record other options for handling signals)



Generating a Signal

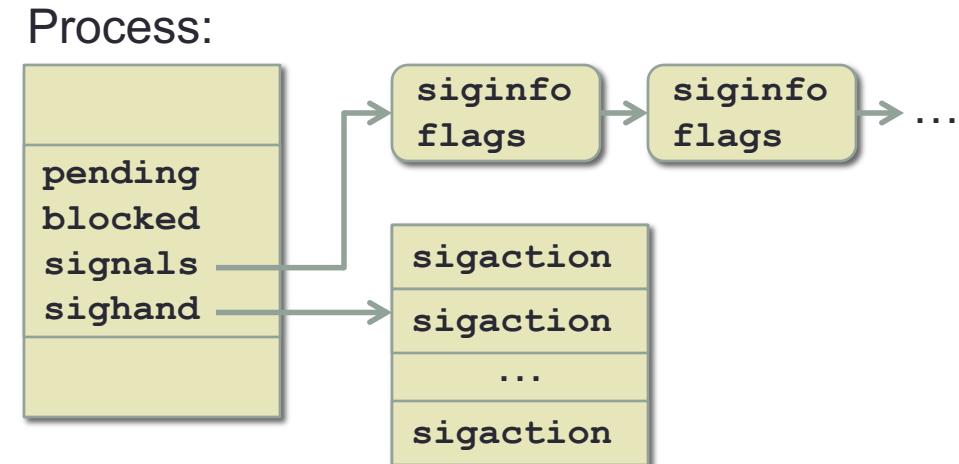
- When a signal is sent to a process:
 - The kernel invokes a specific function to update the process' signal structures, perform scheduling tasks, etc.
 - e.g. Linux 2.6 has `specific_send_sig_info()` kernel function
- If the process already has a pending signal of that type, the new signal is ignored
 - For real-time signals, this test is skipped; every occurrence of a real-time signal is delivered
- If the process is ignoring the signal, nothing is done
 - No structures are updated
 - No scheduling tasks occur

Process:



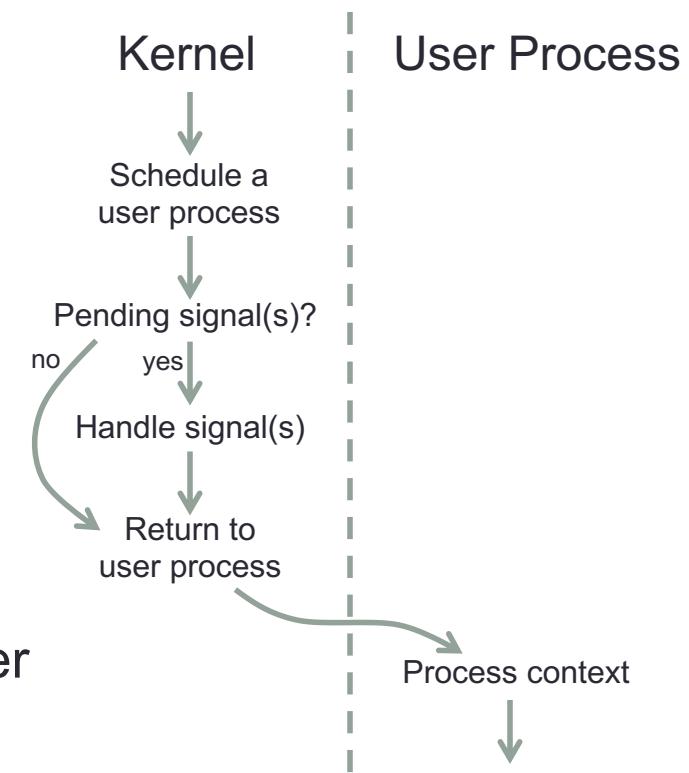
Generating a Signal (2)

- Otherwise, new signal is appended to the signal queue
 - The **pending** bit-vector is also updated
- If the process is currently blocked or suspended, it is moved to the ready state
- Note: a few signal types are not added to signal queue
 - e.g. **SIGKILL**, **SIGSTOP**
 - These signals are enforced immediately by the kernel the next time the process runs
 - Affects the process' execution state in the scheduler



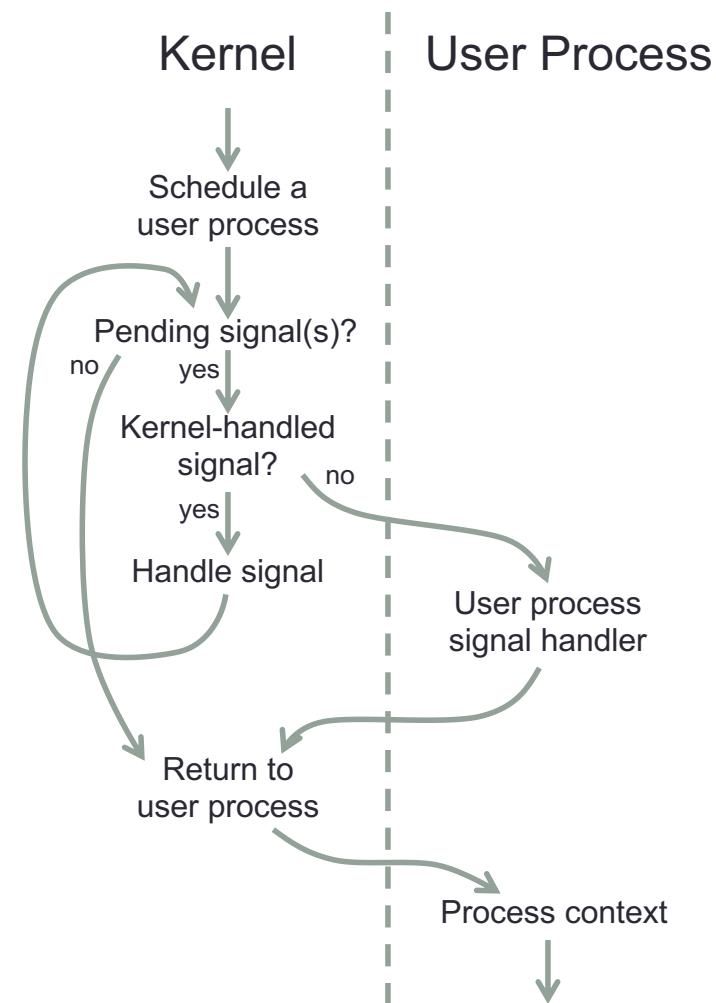
Delivering Signals

- Signals are only delivered to the currently running process
 - Obvious; the process must hold the CPU to run the signal handler
- The kernel checks for pending signals when it is about to return back to the user process
 - Kernel checks the process' signal state
 - If there are pending signals to deliver, they are delivered at this point
- Two ways a signal can be handled:
- Signal is ignored, or default action is to be performed
 - These signals are handled by the kernel
- Signal has a user-mode handler
 - Kernel must invoke the user-mode handler



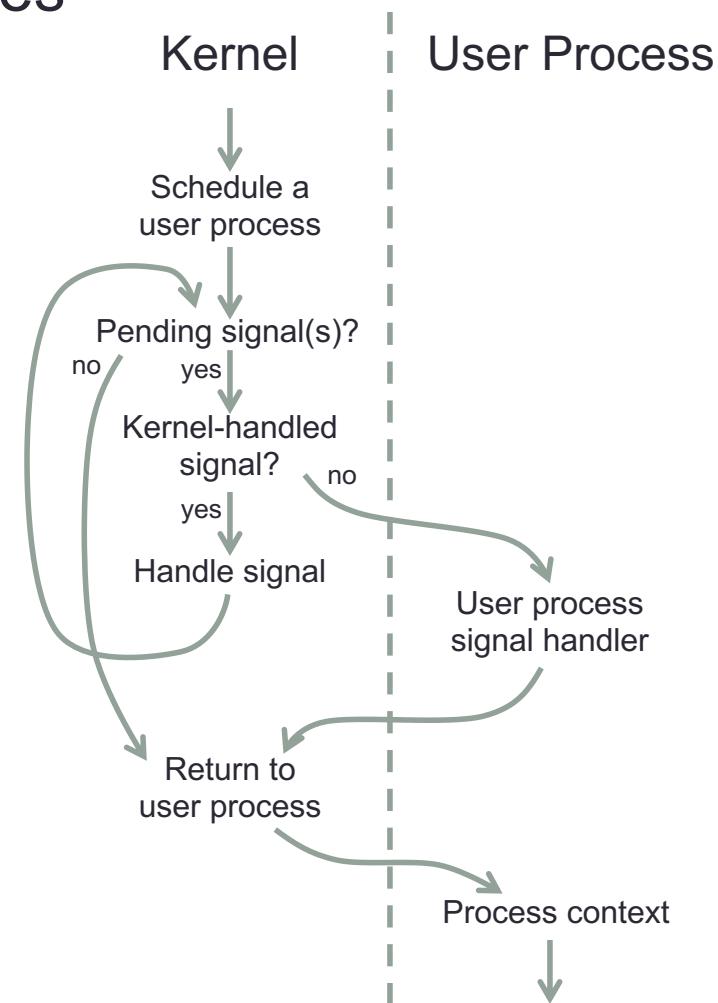
Delivering Signals (2)

- Two ways a signal can be handled:
 - Signal is ignored, or the default action is to be performed
 - Signal has a user-mode handler
- For kernel-handled signals, it can handle as many as are pending
- For user-process-handled signals, only one signal is handled
 - Other pending signals will be delivered the next time the scheduler is invoked
- In Linux 2.6 kernel, signal delivery handled by `do_signal()` function
 - This code has been greatly restructured in subsequent Linux kernel releases



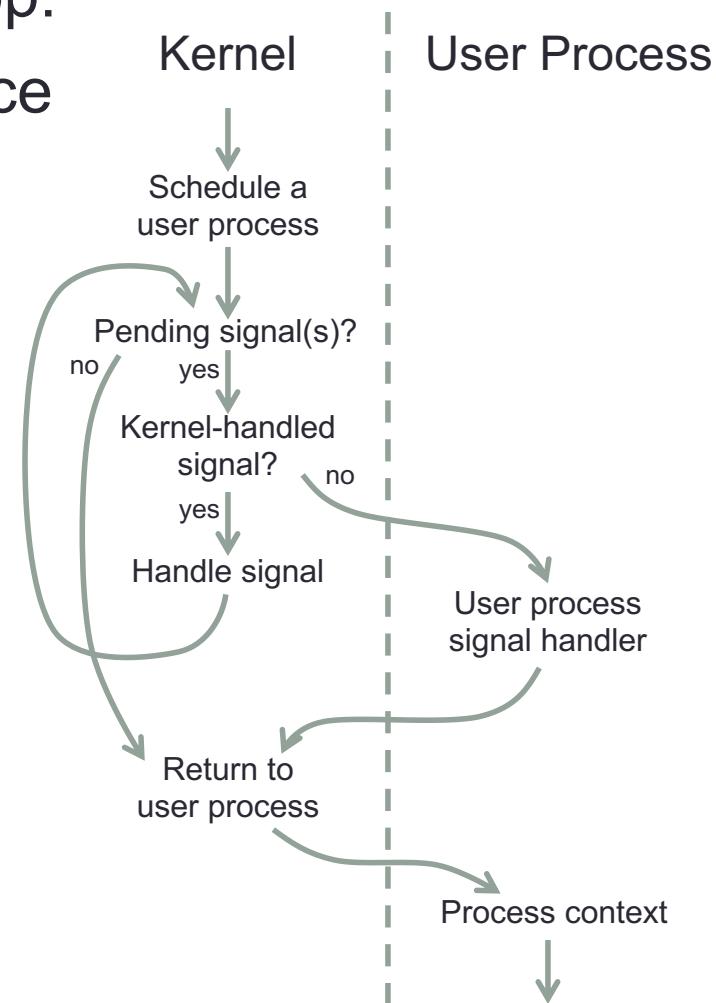
Delivering Signals (3)

- Several big signal delivery challenges
- Signal handler must return back to the kernel so that the previous user process context can be restored
- Signals can interrupt system calls
 - Particularly on preemptible kernels
- Handlers can make system calls
 - When returning from syscall handler, must return to the signal handler, not the interrupted user process context
- Signals can `siglongjmp()` from the signal handler back to another part of the user process



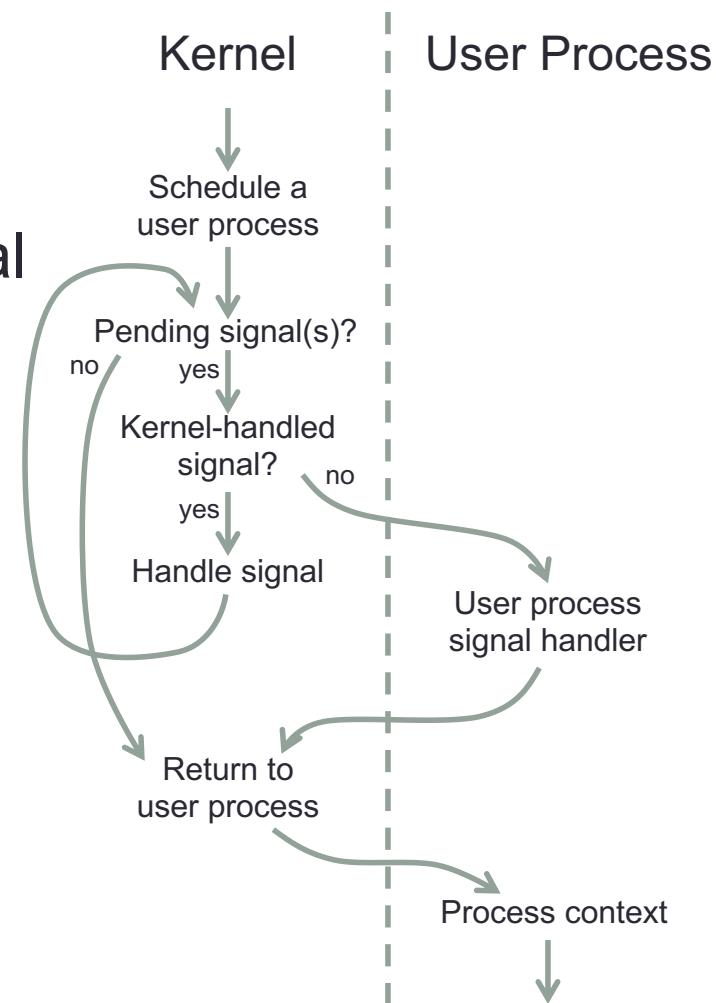
Delivering Signals (4)

- Linux 2.6 `do_signal()` uses a loop:
- Try to find a pending signal to service
 - If there are multiple pending signals, kernel may choose to service them in a different order from queue order
 - (Some signals also cancel each other, like `SIGSTOP` and `SIGCONT`)
- If no more pending signals, the kernel returns to the user process
- If a pending signal was found, it is removed from pending queue
 - User process state is also updated to mark the signal as no longer pending



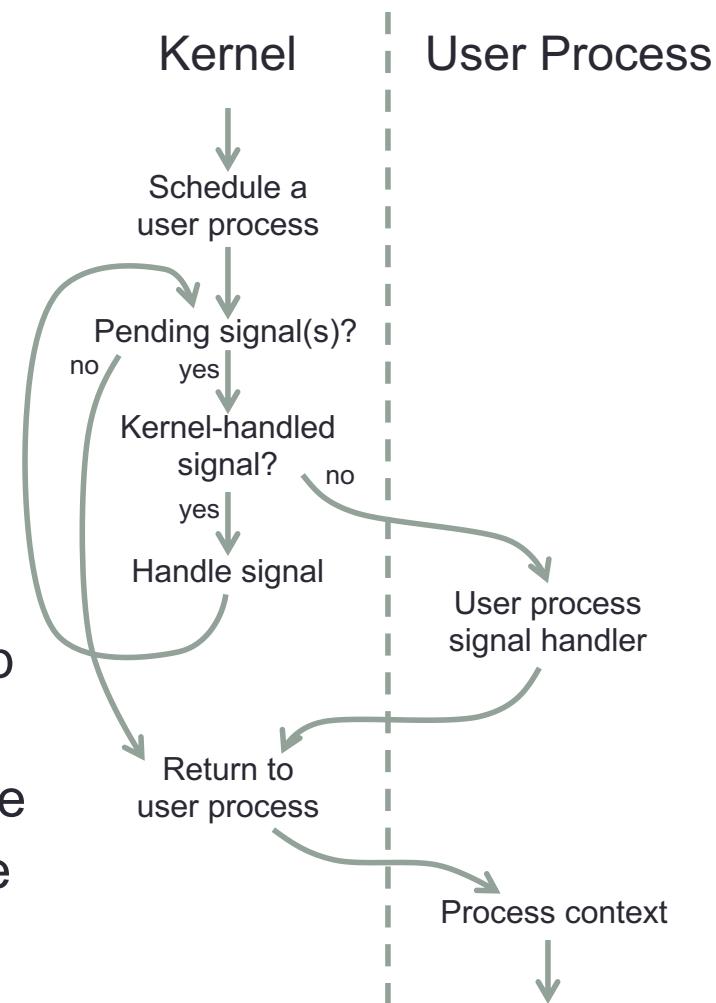
Delivering Signals (5)

- Linux 2.6 `do_signal()`, cont.
- If signal is currently being ignored, `do_signal` continues to next signal
 - A signal may be ignored if the process has set its handler to `SIG_IGN`
 - Or, if signal's handler is `SIG_DFL` and default action is to ignore the signal, the signal is ignored
 - e.g. `SIGCHLD` is ignored by default



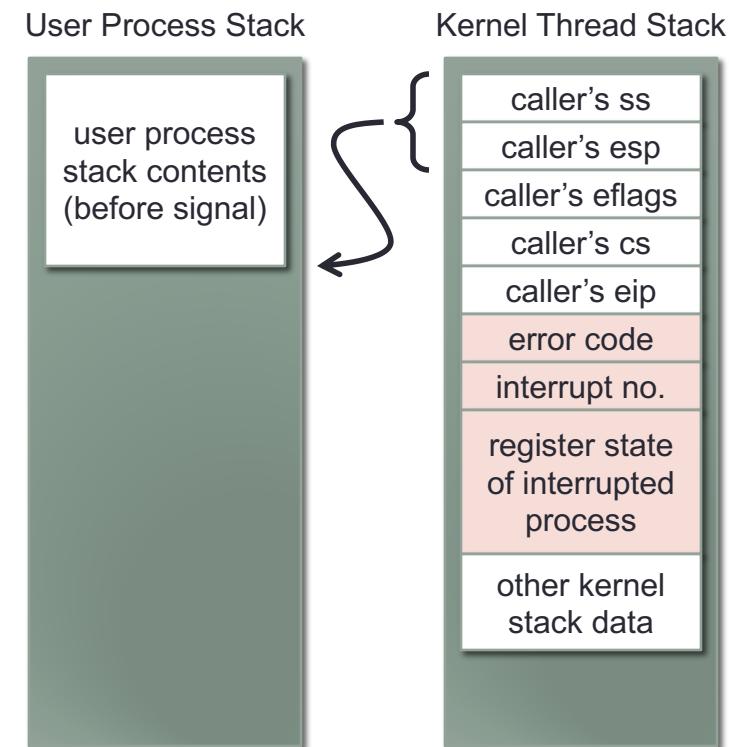
Delivering Signals (6)

- Linux 2.6 `do_signal()`, cont.
- Next, `do_signal` checks if the default action should be used
 - Signal's handler is set to `SIG_DFL` (and default action is not to ignore)
 - If so, `do_signal` carries out the action
- Default actions:
 - **Terminate** – kill the process
 - **Dump** – kill process, create a core dump
 - **Ignore** – ignore signal (handled earlier)
 - **Stop** – move process to suspended state
 - **Continue** – move process to ready state



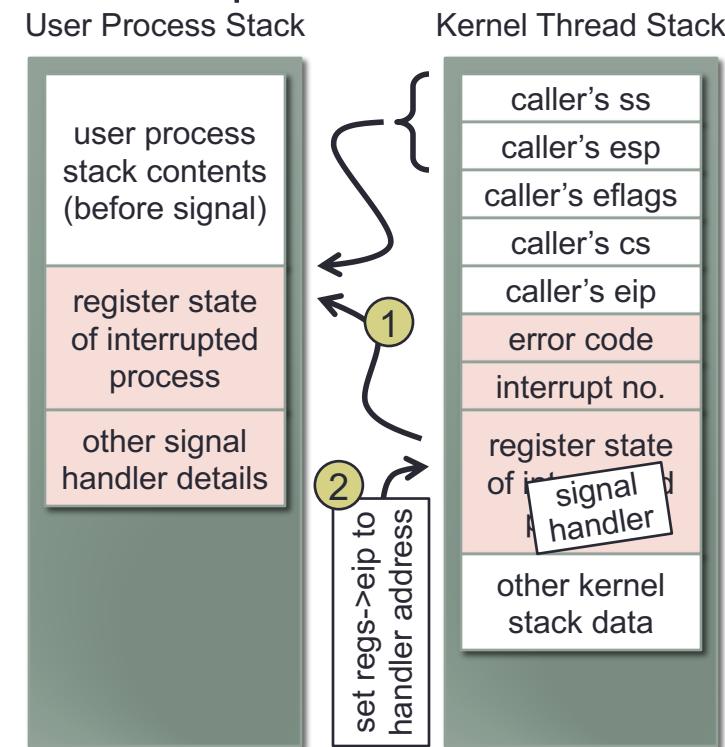
Delivering Signals (7)

- If none of the previous cases hold, `do_signal` must invoke the user process' signal handler
- Recall the state of our stacks within `do_signal` call:
- It's easy to cause the user process to run the signal handler...
 - Just set the process' EIP to the address of the signal handler
- Problem: can't just overwrite the previous CPU state of the process
 - When signal handler completes, must return to whatever was interrupted in the user process
- Must set up a new CPU context for the signal handler to use



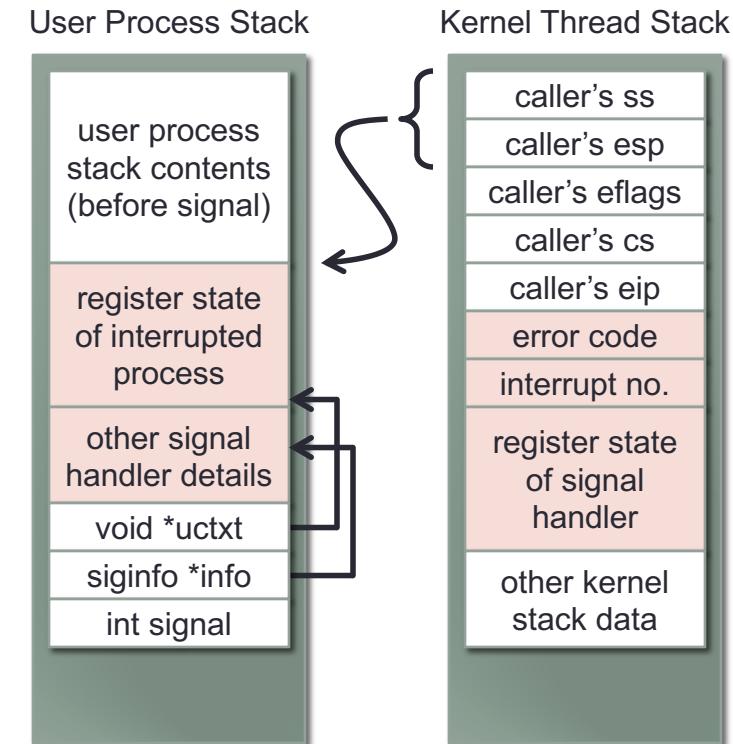
Delivering Signals (8)

- Another problem: the CPU context of the user process' interrupted execution is on the kernel stack...
 - ...but kernel stack will be emptied when kernel returns to user mode!
- Solution: **do_signal** copies some critical details to user stack
 - CPU context of the user process before it was interrupted
 - Other details necessary for properly completing signal handler invocation
 - e.g. a bit-vector of blocked signals, other saved registers not in CPU context, etc.
- These details are used when the signal handler returns:
 - The kernel uses them to go back to the previous point in the interrupted process
- Then, the kernel can change the CPU context on the kernel stack to invoke the signal handler



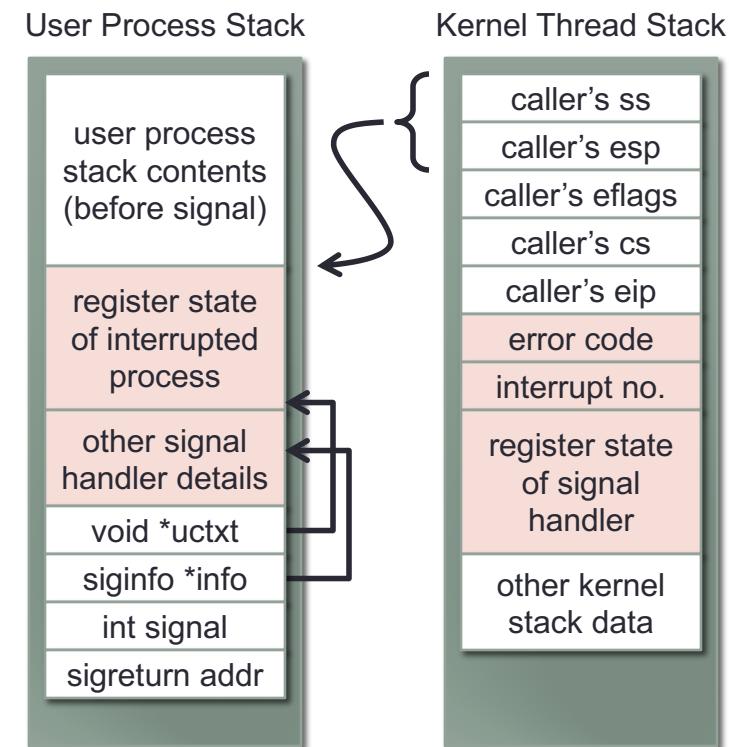
Delivering Signals (9)

- Finally, `do_signal` must set up the stack frame for signal handler function to use
 - If calling a 1-arg signal handler, just push signal # onto stack
 - If calling a 3-arg signal handler, also push signal details onto stack
- Since the process' interrupted CPU context is already on the stack, the “machine context” pointer is easy
 - And, it allows signal handlers to modify the process' CPU state directly, e.g. to implement user-mode threading libraries



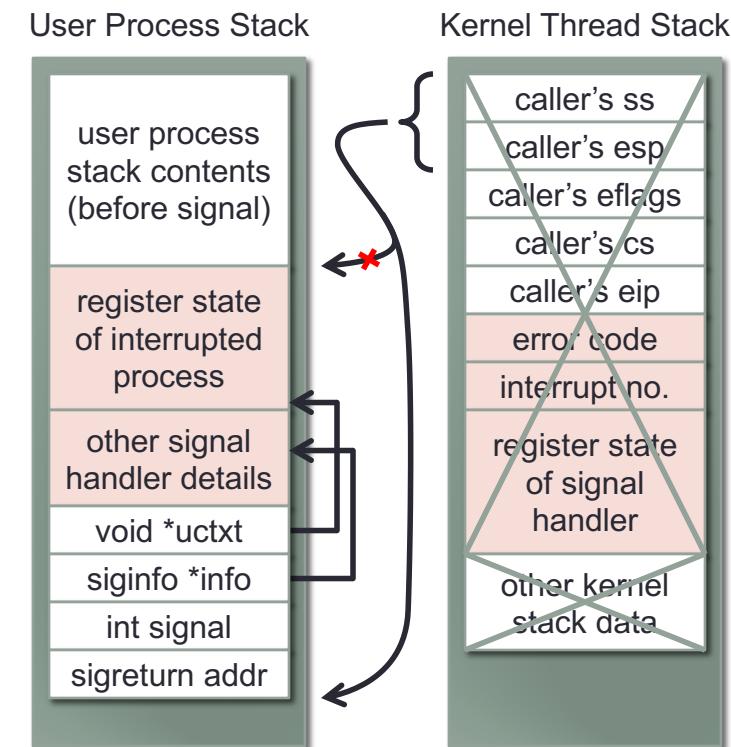
Delivering Signals (10)

- One last component needed for the stack frame:
 - The return address for when the signal handler returns
- Need the signal handler to return to the kernel:
 - Allow kernel to complete final signal-handling tasks, and restore the interrupted process' original context
- The kernel inserts address of code to invoke the **sigreturn** syscall
 - i.e. a wrapper to code that executes “`mov NR_sysreturn, %eax; int $0x80`”
- **sigreturn** has a single purpose:
 - Perform the final task of restoring the interrupted process' CPU context from its location on the stack



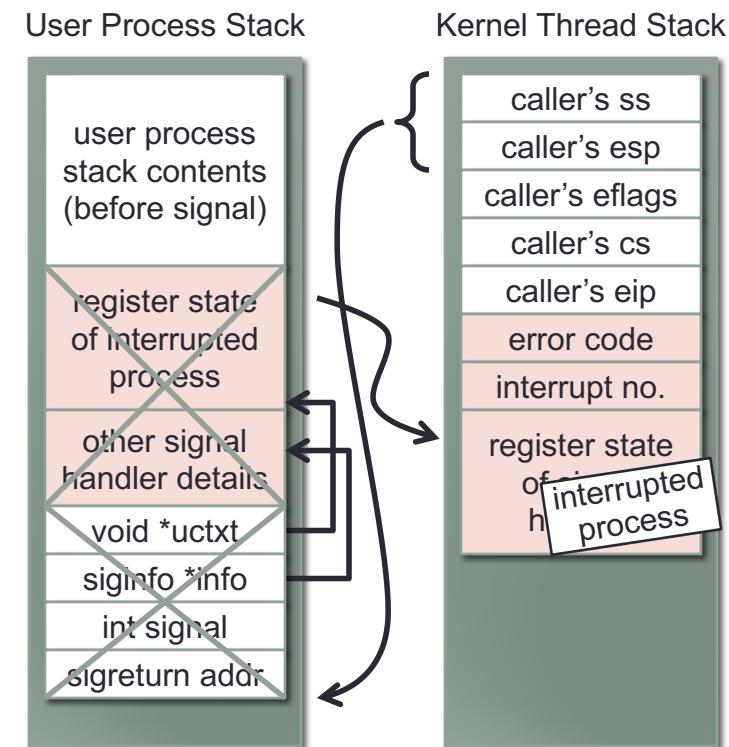
Delivering Signals (11)

- Since user stack was modified, must update caller's `esp`
- Finally, `do_signal()` is done
 - `do_signal` returns to its caller inside the kernel...
 - The kernel returns to user mode...
 - Register state of the user process is restored from the kernel stack...
 - User process begins executing the signal handler!
- Signal handler can make system calls with no problems
 - Original CPU context of the interrupted process is safely stored on user stack



Delivering Signals (12)

- When the signal handler returns, the **sysreturn** system call is invoked
 - Note: again the kernel thread stack contains user-process details
- The **sysreturn** syscall copies the original CPU context of interrupted process back into the kernel stack
 - (along with any changes to the CPU context made by the signal handler)
- Then, **sysreturn** returns back to the user process
 - Resumes process execution at the point where the signal interrupted the process
 - User stack pointer is restored as well, eliminating now-unneeded stack data

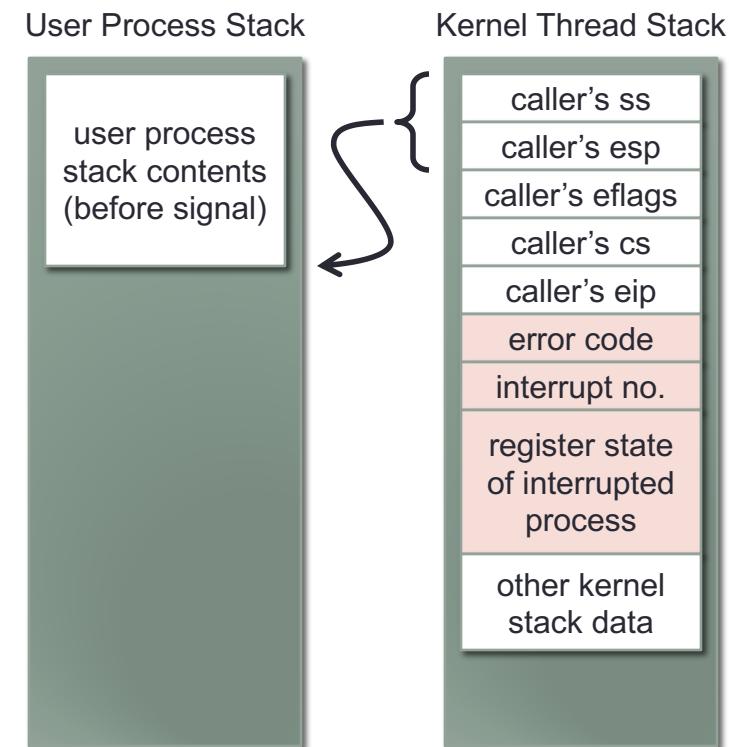


Signals and System Calls

- Process may be blocked on a syscall when signal occurs
 - The system call is interrupted by the signal
 - Generally, some action must be taken by the kernel in this case
 - e.g. the system call might return `EINTR` as its error code
- Some system calls can be automatically restarted when the signal occurs
 - e.g. `read()` or `write()` may be automatically restarted if they hadn't performed any work by the time the signal occurs
- Some system calls must report that they were interrupted
 - e.g. if `nanosleep()` is interrupted by a signal, it returns `EINTR`, and reports the amount of time remaining in the sleep interval

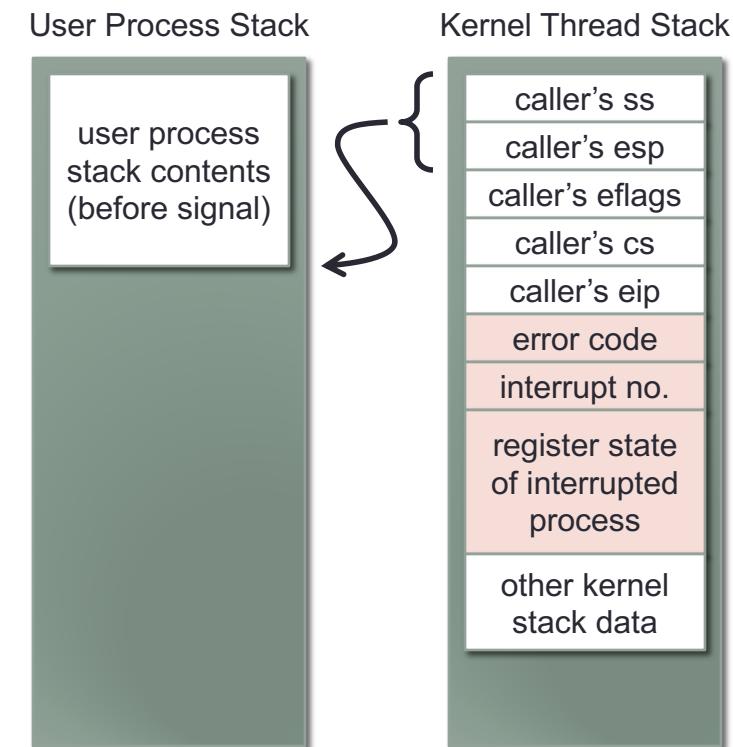
Signals and System Calls (2)

- If a process is in a system call when it receives a signal:
- It was previously in kernel code when signal is delivered
 - It entered the kernel by performing an `int $0x80` trap operation (or a fast system-call, e.g. using `sysenter` instruction)
- Or, the process may be in some blocked queue
- The signal delivery mechanism will see the interrupt # in process' state
 - Original value of `eax` is also stored, indicates which syscall was interrupted



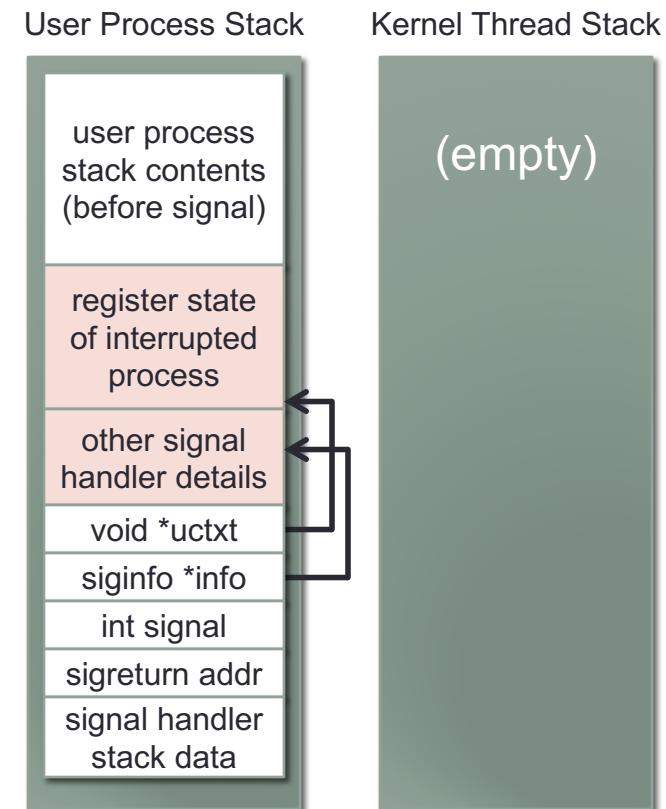
Signals and System Calls (3)

- Depending on the process' configuration, and on system call that was interrupted, kernel takes one of two actions:
- Option 1: the kernel modifies the interrupted process' CPU context to show the system call as interrupted
 - Sets the process' **eax** register to **-EINTR**
 - The process will see the system call return this error code
- Option 2: the kernel modifies process' CPU context to rerun the system call
 - Sets the process' **eax** register to the original system call number
 - Subtracts 2 from process' **eip** register to force **int \$0x80** (or **sysenter**) to run again



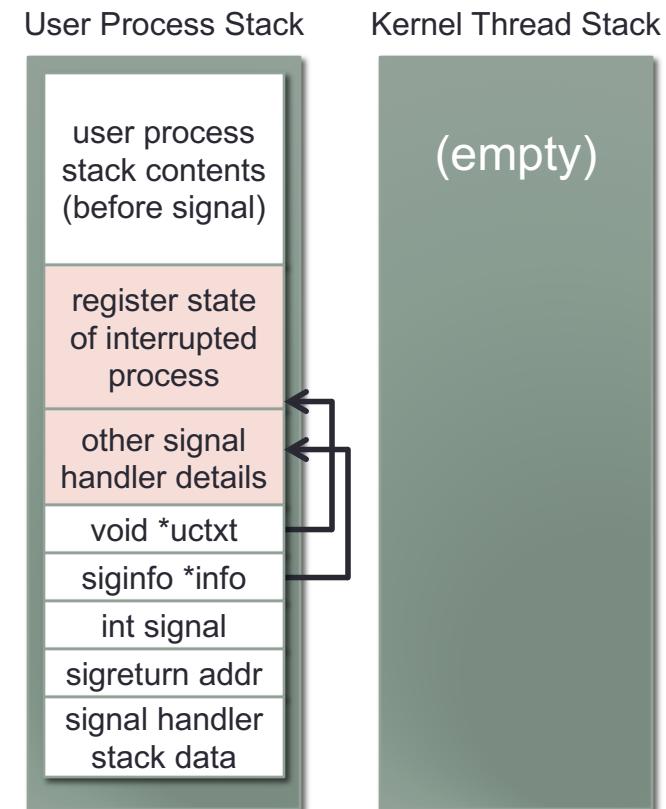
Signal Handlers and `siglongjmp`

- Signal handlers can use `siglongjmp()` to jump to another part of the user program
 - Effectively terminates the signal handler and resumes execution of the user program at a different point
 - Described as a “non-local goto”
- In fact, could use `longjmp()` or `siglongjmp()`
 - `sigsetjmp` / `siglongjmp` are strongly preferred because they can optionally save and restore blocked-signals mask
 - In practice, that is the only difference between `setjmp` and `sigsetjmp`



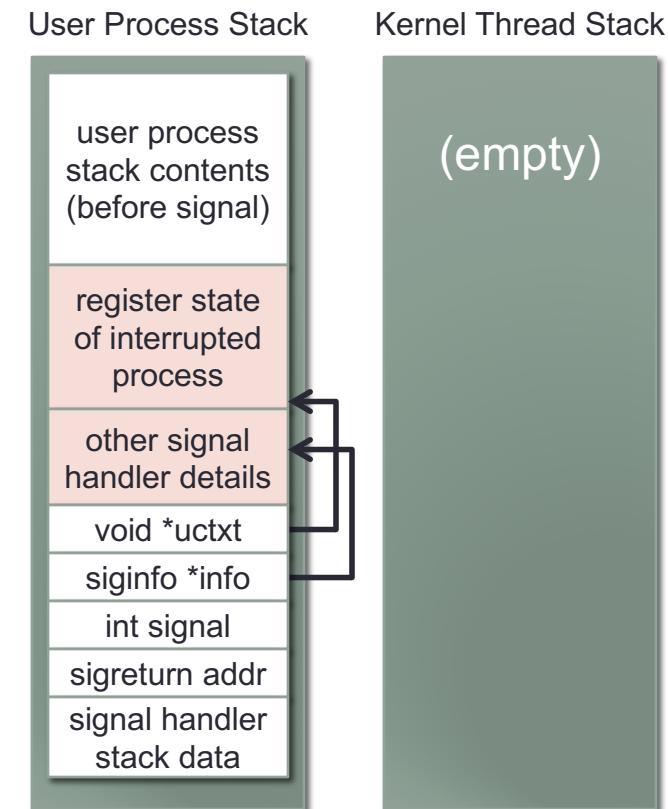
Signal Handlers and `siglongjmp` (2)

- Does performing a `siglongjmp` from a signal handler break the kernel's signal-delivery machinery at all?
- Answer should be obvious from the stack state, and from previous discussion
- Recall:
 - When a pending signal is dequeued by `do_signal()`, process' state is updated to show the signal as no longer pending
 - Kernel doesn't *require* the signal handler to return in order to "complete" processing
 - As far as kernel is concerned, signal is done!
 - Only task is to restore blocked signal mask (which `siglongjmp` can also take care of)



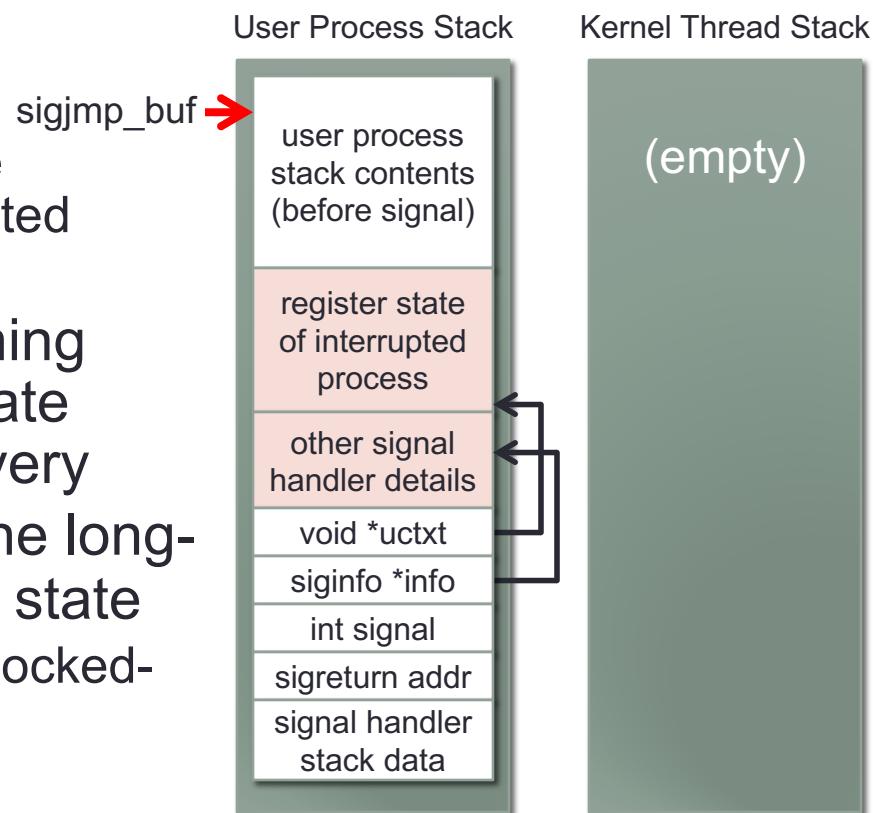
Signal Handlers and `siglongjmp` (3)

- Also: kernel will always check for pending signals to deliver, when returning to the user process
 - Any pending signals will likely be delivered the next time the scheduler starts running the process again
- Finally, the kernel's signal-delivery mechanism leaves no state on the kernel stack
 - All state for both the signal handler and the interrupted point in the process is contained within the user stack
- If a user process long-jumps out of signal handler, it will have no effect on delivery of any future signals



Signal Handlers and `siglongjmp` (4)

- Why would a process long-jump out of signal handler?
 - Needs to have a jump-buffer that records the CPU state at some earlier point in stack
- A process that long-jumps out of a signal handler *doesn't want to* restore the old execution state!
 - The process *doesn't want to* resume executing the code that was interrupted when the signal occurred
- Long-jump will discard all intervening execution state, including CPU state saved by the kernel at signal delivery
- No need to invoke `sysreturn`; the long-jump will restore the desired CPU state
 - And, `siglongjmp` will restore the blocked-signal mask to an appropriate state



Next Time

- Kernels maintain complex structures on behalf of user processes...
- Kernel has a very limited amount of memory for dynamic allocation
 - Kernel has a fixed memory area for data that pertains to all processes
 - And allocation needs to be fast
- Next time: kernel allocators

